

Memory Interleaving

- As CPU processing speed increased, it started to outpace memory access speed. This resulted in longer wait times as the CPU idled as the slower memory unit retrieved data.
- One strategy developed to overcome this was memory interleaving. It is a design which compensates for the slow speed of RAM by spreading memory addresses across multiple chips.

Memory Interleaving


- Computer memory consists of a linear array of addressable storage cells that are similar to registers.
- Memory can be **byte-addressable**, or **word-addressable**, where a word typically consists of two or more bytes.
- Memory is constructed of RAM chips, often referred to in terms of length \times width (L \times W).
- A **4M** \times **8** RAM chip gives us 4,194,304 8-bit memory locations.

Memory Organization

- How does the computer access a memory location corresponding to a particular address?
- We observe that 4MB can be expressed as $2^2 \times 2^{20} = 2^{22}$ words.
- The memory locations for this memory are numbered 0 through $2^{22} - 1$.
- Thus, the memory bus of this system requires **22 address lines**.
 - The address lines “count” from 0 to $2^{22} - 1$ in binary. Each line is either “on” or “off” indicating the location of the desired memory element.

Memory Organization

- Physical memory usually consists of more than one RAM chip.
- Access is more efficient when memory is organized into banks of chips with the addresses **interleaved** across the chips.



Memory Organization

- Suppose we need to build 32K x 8, byte-addressable memory, but we only have 2K x 8 RAM chips.
- Connect 16 rows of chips together:

Row 15	2K x 8
Row 14	2K x 8

Row 1	2K x 8
Row 0	2K x 8

Memory Organization

- Each chip addresses 2K bytes.
- Because we have a total of 32K addresses for this memory, must have 15 bits ($2^5 \times 2^{10} = 2^{15}$ bytes to access).
- Each chip only holds 2^{11} bytes, so 4 bits are used to determine the address of each chip (16 of them = 2^4).

Memory Organization

- Either the leftmost bits or the rightmost bits are used, depending on the architecture.
- A decoder is used to decode these 4 bits to determine which chip holds the desired data.
- Once the proper chip has been located, the remaining 11 bits are used to determine the **offset** on the chip (where the data is).

Memory Organization

- A single memory module can only be accessed sequentially so memory is normally split across multiple modules (or banks) which can be accessed simultaneously.
- This process is called **Memory Interleaving**.
- With **low-order interleaving**, the low-order bits are used to address the banks.
- With **high-order interleaving**, the high-order bits are used to address the banks.

Memory Organization

- Suppose we have a byte-addressable memory consisting of 8 modules of 4 bytes each, giving a total of 32 bytes of memory.
- We need 5 bits to uniquely identify each byte ($32 = 2^5$) in the total memory space.
- We have 8 modules so we need 3 bits to address each of them ($2^3 = 8$).
- The remaining 2 bits are used to determine the offset within the module.

Memory Organization

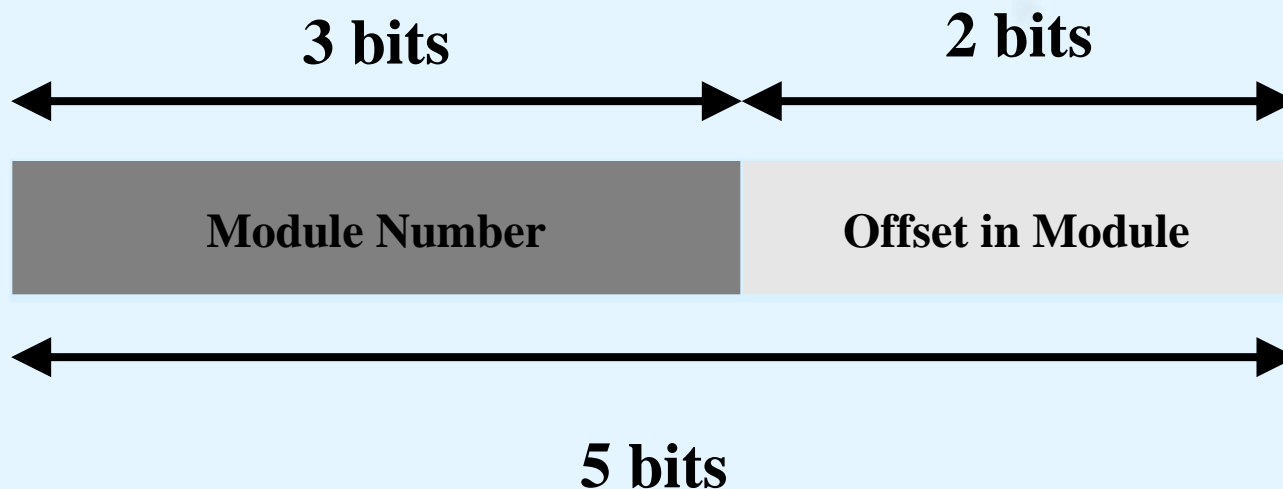
- **High-order interleaving** distributes the addresses so that each module contains consecutive addresses.
- Module 0 contains the data stored at addresses 0, 1, 2 and 3
- Module 1 contains the data stored at addresses 4, 5, 6 and 7.
-
- Module 7 contains the data stored at addresses 28, 29, 30 and 31

Module 0	Module 1	Module 2	Module 3	Module 4	Module 5	Module 6	Module 7
0	4	8	12	16	20	24	28
1	5	9	13	17	21	25	29
2	6	10	14	18	22	26	30
3	7	11	15	19	23	27	31

High-Order Interleaving

Memory Organization

- High-order interleaving uses the first 3 bits to determine the address of the memory module and the remaining 2 bits are used to determine the offset within the module.



Memory Organization

Module	Decimal Address	Binary Address	Split	Module Number	Offset
Module 0	0	00000	000 00	0	0
	1	00001	000 01	0	1
	2	00010	000 10	0	2
	3	00011	000 11	0	3
Module 1	4	00100	001 00	1	0
	5	00101	001 01	1	1
	6	00110	001 10	1	2
	7	00111	001 11	1	3

Memory Organization

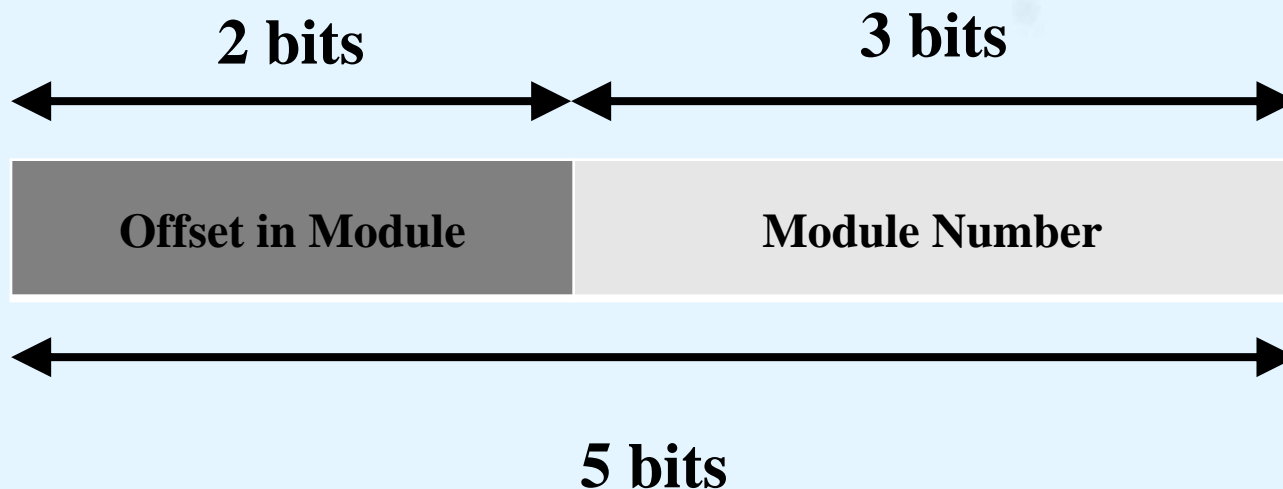
- **Low-order interleaving** places consecutive memory addresses in different modules
- Module 0 contains the data stored at addresses 0, 8, 16 and 24
- Module 1 contains the data stored at addresses 1, 9, 17 and 25.
-
- Module 7 contains the data stored at addresses 7, 15, 23 and 31

Module 0	Module 1	Module 2	Module 3	Module 4	Module 5	Module 6	Module 7
0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31

Low-Order Interleaving

Memory Organization

- Low-order interleaving uses the last 3 bits to determine the address of the memory module and the remaining 2 bits are used to determine the offset within the module.



Memory Organization

Module	Decimal Address	Binary Address	Split	Offset	Module Number
Module 0	0	00000	00 000	0	0
	8	01000	01 000	1	0
	16	10000	10 000	2	0
	24	11000	11 000	3	0
Module 1	1	00001	00 001	0	1
	9	01001	01 001	1	1
	17	10001	10 001	2	1
	25	11001	11 001	3	1

Memory Organization

- The order of interleaving refers to the number of memory banks used in the memory:
 - 4-way uses 4 banks
 - 8-way uses 8 banks
 - 16-way uses 16 banks, etc.
- The number of bits needed to identify the banks is k where $2^k =$ the number of memory banks.
- So, for 8-way interleaving, $k=3$ as $2^3 = 8$, so we need 3 bits.
- For 16-way interleaving, we need 4 bits.

Memory Organization

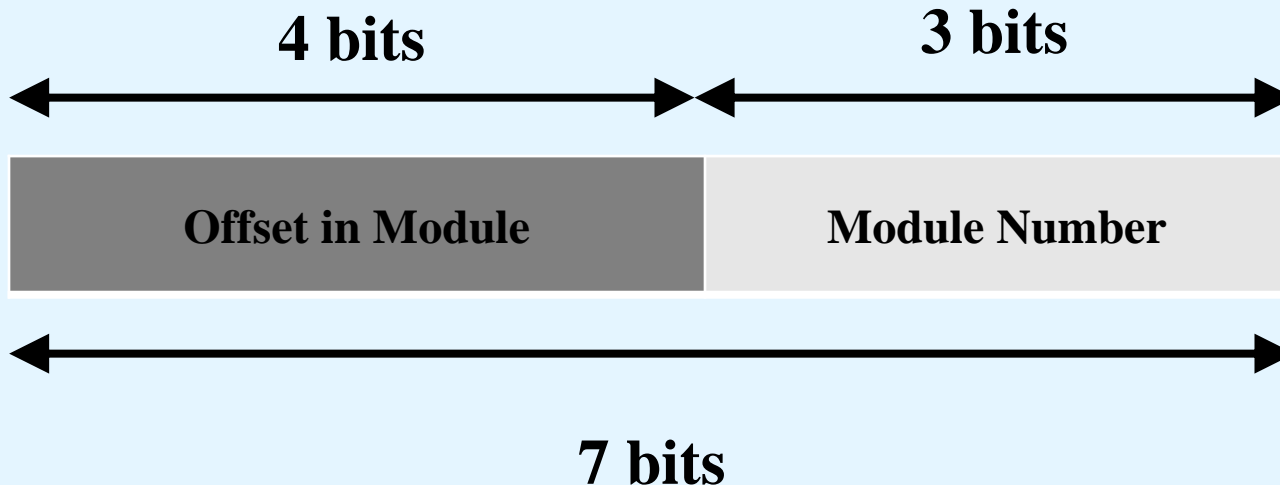
- So for n-way interleaving, we need $n = 2^k$ bits for addressing the memory banks.

EXAMPLE:

Suppose we have a 128 byte memory and we are using **8-way, low-order interleaving**. What is the structure of the addresses for this memory?

Memory Organization

- We have 128 bytes, so we need 7 bits for each individual address ($128 = 2^7$).
- The interleaving is low order, so the lowest-order bits are used for the memory bank address.
- We are using 8-way interleaving, so we have 8 memory banks (modules), so we need 3 bits to address all of them ($8 = 2^3$).




Memory Organization

- Note that each module must be of size 2^4 .
- This can be concluded in two ways:
 1. We have 128 bytes and we have 8 modules, so each module will hold 16 bytes: $(128/8 = 2^7/2^3 = 2^4 = 16)$.
 2. We can also see from the address structure that the offset in the module requires 4 bits, allowing for $2^4 = 16$ bytes per module.


High-Order Interleave

- Good if Modules can be accessed independently by different units, e.g. by the CPU and a Hard Disk (or a second CPU) **AND the units use different Modules**
- => Parallel operation => Higher Performance



Low-Order Interleave

- Good if the CPU (or other unit) can request multiple adjacent memory locations.



- Since adjacent memory locations lie in different Modules an “advanced” memory system can perform the accesses in parallel. Such adjacent accesses often occur in practice, e.g.
 - i) Elements in an array, e.g. `Array[N]`, `Array[N+1]`, `Array[N+2]`,
 - ii) Instructions in a Programs, `InstructionN`, `InstructionN+1`,...
- In the above situations, an “advanced” CPU can pre-fetch the adjacent memory locations => higher performance.

Memory Organization

Exercise 1:

For a 32K x 8 memory that uses **16-way, high-order interleaving**, find the location (chip and offset) of the following address:

001000000100111

Memory Organization

Answer:

32K memory by 16-way = 16 x 2K chips.

$32K = 2^5 \times 2^{10} = 2^{15} = 15$ bits for all addresses.

$16 = 2^4 = 4$ bits for chip addresses.

$15 - 4 = 11$ bits for offset in each chip.

0010 00000100111

Chip: 2, Offset: 39.

Exercise 2: Repeat for 8-way, low-order interleaving.

Memory Organization

Answer:

32K memory by 8-way = 8 x 4K chips.

$32K = 2^5 \times 2^{10} = 2^{15} = 15$ bits for all addresses.

$18 = 2^3 = 3$ bits for chip addresses.

$15 - 3 = 12$ bits for offset in each chip.

001000000100 111

Chip: 7, Offset: 516.